

System Software Crash Course

Samsung Research Russia
Moscow 2019

Block C Compiler Construction
9. Source Code Optimization
Eugene Zouev

Program Optimization

Some general points

- Optimization can be performed on each stage of the program lifecycle: not only while compilation but while design, development and maintenance.
- Do we really need optimization?
The best way to optimize a program - is to design it correctly (then perhaps we do not need to optimize it 😊)
- "Optimization-in-the-small" vs "optimization-in-the-large" ...

Program Optimization

- Finding places in programs which could be optimized (by some criteria) is very much empirical job; in the best case, there is just a set of techniques taken from experience.
- At the same time, there is a number of formal and/or constructive approaches for some kind of optimizations.

Today, we will be discussing **what** to optimize, but not **how** to do this...

Program Optimization

- While source code processing (lexical & syntax analysis)

Big spectrum of optimization techniques.

- While semantic analysis (AST processing).

Sequential AST traversing.

Optimizations depend on the language semantics heavily.

- While target code generation (machine-dependent optimizations)

Depend on the target architecture & on the instruction set.

- While linking: **global** code optimizations.

Example: - C++ code bloat removing.

Elimination of repeated calculations (1)

```
double a = x*(1-sin(y));  
double b = x + y/z;  
double c = y/z + 1 - sin(y);
```



```
double tmp1 = 1-sin(y);  
double tmp2 = z/y;  
  
double a = x*tmp1;  
double b = x + tmp2;  
double c = tmp2 + tmp1;
```

The place:
While AST analysis.

Limitations:

1. Factorized functions cannot issue side effects.
2. Operands of factorized expressions cannot modify their values.

Elimination of repeated calculations (2)

```
static int x = 0;  
double F(double y)  
{  
    x++;  
    return <expression>;  
}
```

Side effect



```
double a = x*(1-F(y));  
double b = x + y/z;  
...  
z = <expression>;  
...  
double c = z/y + 1 - F(y);
```

Modifying value



Elimination of repeated calculations (3)

An expression may look different but still calculate the same value as some other expression => it can also get optimized.

```
double a = b*c - d;
```

```
double e = b;
```

```
b = b + 1 - b*c;
```

```
double f = b*c + c*e;
```

$b*c$: the second calculation of the same value.

The second calculation of $b*c$ results in a different value.

The value of $c*e$ is the same as $b*c$.

Replacing slow instructions for faster ones (1)

Actions: comparative performance

Multiplication/division on a power of two

Addition/subtraction

Multiplication

Division


Calculation of an integer power

Calculation of an arbitrary power

⇒ Replacing slower operations for faster ones (where possible)

For some target architectures it's **mandatory**: e.g., some RISCs just do not support multiplication!

Replacing slow instructions for faster ones (2)



In general case it's impossible...

- At least one operand must be an integer constant.
- The constant should be relatively small; otherwise rounding errors will accumulate.

Replacing slow instructions for faster ones (3)

Division

X/K



Multiplication

$X * (1/K)$

K - constant

```
double x = c/b;  
double y = (e+f)/b + d;  
double z = b;  
b = b+1;  
...  
z = sin(x)/z + e/b;
```



```
double tmp = (double)1/b;  
double x = c*tmp;  
double y = (e+f)*tmp + d;  
double z = b;  
b = b+1;  
...  
z = sin(x)*tmp + e/b;
```

Excluding redundant calculations

```
double a;  
...  
a = (x+y)*sin(z);  
...  
a = x/y;
```



```
double a;  
...  
a = x/y;
```

If the value of **a** does not change between two assignments, then the first assignment can be removed.

Limitation: the action being removed cannot make side effects.

Constant propagation (1)

If the value of a variable is known
then the variable reference could be
replaced for the value itself.

```
double a = 2.0;  
double b = 3.5;
```

```
...  
double c = a*b;
```

```
...  
double t = (b+c)*a+x;
```



```
double a = 2.0;  
double b = 3.5;
```


```
...  
double c = 7.0; // a*b
```

```
...  
double t = 10.5 + x; // (b+c)*a
```

Constant propagation (2)


If the value in a loop condition is known in advance then the loop could be simplified.

```
while ( <Expression> )  
{  
    <Statements>  
}
```



```
Loop:  
    if ( !<Expression> ) goto Exit;  
    <Statements>  
    goto Loop;  
Exit:  
    ;
```

```
while ( 1 )  
{  
    <Statements>  
}
```



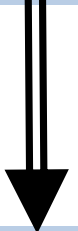
```
Loop:  
    if ( !<Expression> ) goto Exit;  
    <Statements>  
    goto Loop;  
Exit:  
    ÷
```

Type conversion optimizations

Type conversion is a potentially costly operation; therefore it's a good candidate for optimizations.

```
double a, b;  
long i, j;  
...  
double c = a + i + b - j;
```

... a + (double)i + b - (double)j ...



```
double a, b;  
long i, j;  
...  
double c = (a+b) + (i-j);
```

... a+b + (double)(i-j) ...

Reducing calls & indexings

Access to array elements and function calls are also good candidates for optimizations...

```
for i:integer range 1..100 loop
  x(i) = y(i)+1/y(i);
  z(i) = y(i)**2;
end loop;
```

Address of **y(i)** gets
calculated **300 times**



```
for i:integer range 1..100 loop
  declare
    tmp : real := y(i);
  begin
    x(i) = tmp+1/tmp;
    z(i) = tmp**2;
  end;
end loop;
```


Address of **y(i)** gets
calculated **100 times**

Merging loops (1)

Loops are main consumers of CPU time!

```
for i:integer range 1..100 loop  
  x(i) = 0;  
end loop;
```

```
for i:integer range 1..100 loop  
  z(i) = y(i)**2;  
end loop;
```



```
for i:integer range 1..100 loop  
  x(i) = 0;  
  z(i) = y(i)**2;  
end loop;
```

Costs for loop organization
are reduced

Merging loops (2)

```
for i:integer range 1..100 loop  
  x(i) = 0;  
end loop;
```

```
for i:integer range 1..200 loop  
  z(i) = y(i)**2;  
end loop;
```



```
for i:integer range 1..100 loop  
  x(i) = 0;  
  z(i) = y(i)**2;  
end loop;
```

```
for i:integer range 101..200 loop  
  z(i) = y(i)**2;  
end loop;
```


(More general case)

Overall amount
of iterations: 300

Overall amount
of iterations: 200

Unrolling loops (1)

```
for (int i=0; i<100; i++)  
{  
    x[i] = y[i]*z[i];  
}
```



```
for (int i=0; i<100; i+=2)  
{  
    x[i] = y[i]*z[i];  
    x[i+1] = y[i+1]*z[i+1];  
}
```

Loop step = 1
Overall amount
of iterations: **100**

Loop step = 2
Overall amount
of iterations: **50**

Unrolling loops (2)

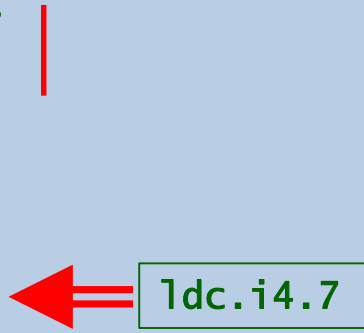
```
-- Skip past blanks, loop is opened up for speed
while Source (Scan_Ptr) = ' ' loop
  if source (Scan_Ptr + 1) /= ' ' then
    Scan_Ptr := Scan_Ptr + 1; exit;
  end if;
  if source (Scan_Ptr + 2) /= ' ' then
    Scan_Ptr := Scan_Ptr + 2; exit;
  end if;
  if source (Scan_Ptr + 3) /= ' ' then
    Scan_Ptr := Scan_Ptr + 3; exit;
  end if;
  if source (Scan_Ptr + 4) /= ' ' then
    Scan_Ptr := Scan_Ptr + 4; exit;
  end if;
  if source (Scan_Ptr + 5) /= ' ' then
    Scan_Ptr := Scan_Ptr + 5; exit;
  end if;
  if source (Scan_Ptr + 6) /= ' ' then
    Scan_Ptr := Scan_Ptr + 6; exit;
  end if;
  if source (Scan_Ptr + 7) /= ' ' then
    Scan_Ptr := Scan_Ptr + 7; exit;
  end if;
  Scan_Ptr := Scan_Ptr + 8;
end loop;
```

A real example:
The scanner of the
Ada GNAT compiler

Stack optimizations (1)

```
class Program
{
    int F(int a,int b)
    {
        int c = 7;
        int x = (a-b)*(a+c);
        return x;
    }
}
```

```
...
IL_0001: ldc.i4.7
IL_0002: stloc.0
IL_0003: ldarg.1
IL_0004: ldarg.2
IL_0005: sub
IL_0006: ldarg.1
IL_0007: ldloc.0
IL_0008: add
IL_0009: mul
IL_000a: stloc.1
...
```



Even such a simple code is **not optimal**;
it could be improved.

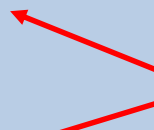
The first improvement:
constant propagation

Stack optimizations (2)

```
class Program
{
    int F(int a,int b)
    {
        int c = 7;
        int x = (a-b)*(a+c);
        return x;
    }
}
```

```
...
IL_0001: ldc.i4.7
IL_0002: stloc.0
IL_0003: ldarg.1
IL_0004: ldarg.2
IL_0005: sub
IL_0006: ldarg.1
IL_0007: ldloc.0
IL_0008: add
IL_0009: mul
IL_000a: stloc.1
...
```

The value of **a**
gets pushed
two times



```
...
IL_0003: ldarg.1
           dup
IL_0004: ldarg.2
IL_0005: sub
           swap
IL_0006: ldarg.1
IL_0007: ldloc.0
IL_0008: add
IL_0009: mul
IL_000a: stloc.1
...
```

The second improvement:

Replace the second push instruction for faster instruction(s).

dup Copies the value on the top of the stack

swap Swaps positions of two stack topmost operands
(exists in JVM, but **doesn't** in MSIL)

VM: stack- or register-based?

```
VAR
  ch: CHAR;
  lint: LONGINT;
BEGIN
  lint := LONG(ORD(ch));
  ...
```

```
MOVZBD ch, lint
  // move, zero-extending
  // byte to double-word
```

National Semiconductor 32000 family

```
LOAD1 ch    // move single_byte CHAR value onto stack
ORD          // zero-extend, result is an INTEGER
LONG         // sign-extend, result is a LONGINT
STORE4 lint // move four_byte LONGINT value back to memory
```

An abstract stack computer

Michael Franz
Code-Generation On-the-Fly: A Key to Portable Software
Dissertation ETH No. 10947, 1994

Tail Recursion (1)

```
void f(int x)
{
    if ( x == 0 ) return;
    ...Some actions...
    f(x-1);
}
```

```
void f(int x)
{
    if ( x == 0 ) return;
    ... Some actions...
    if ( Some_condition )
        f(x-1);
    else
        f(x-2);
}
```

The idea:

If the recursive call is the very last operation in the function body then it could be replaced for the direct jump to the beginning of the body - perhaps with argument (re)initialization.

See more details in:

http://en.wikipedia.org/wiki/Tail_call

Tail Recursion (2)

```
long factorial(long n)
{
    if (n == 0) return 1;
    else return n*factorial(n-1);
}
```

This is **not** tail recursion.
Why?

```
int fac_times(int n, int acc)
{
    if (n == 0) return acc;
    else return fac_times(n-1, acc*n);
}

int factorial(int n)
{
    return fac_times(n, 1);
}
```

Equivalent program **with** tail recursion.