

Compiler Construction: Practical Introduction

System Course for SRR Engineers

Samsung Research Russia
Moscow 2019

Lecture 6


Back End Essentials

- **Compiler Backend**
- **Optimizations for sequential uniprocessor**
- **Static Single Assignment form**
- **Optimizations of parallel computations**

Compiler Backend

Role and Responsibilities

Compiler



Compiler Backend

Responsibility

- Optimally map given program to hardware without compromising correctness

Where is compiler backend?

- Regular compilers (gcc GIMPLE, LLVM)
- Virtual Machines for managed languages (JVM, .NET)
- Scripting languages (Python, JavaScript &.c)
 - Even in browser
- Optimizing execution of a Machine Learning model (e.g. neural network)

Tasks and Issues

- Backend should:
 - Select suitable processor instructions
 - Allocate spare hardware resources such as registers
 - Optimize execution sequence
 - Schedule operations to minimize total execution time
 - &c.
- Backend need to know:
 - How program is executed (Control flow analysis)
 - What values are used during execution (Data flow analysis)
 - “Understand” program’s semantics

Optimizations Examples

- **Conditional Constant Propagation**
 - Constant propagation is a transformation that, given an assignment $x = c$ for a variable x and a constant c , replaces later uses of x with uses of c as long as intervening assignments have not changed the value of x .
- **Tail-recursion optimization**
 - A call from procedure $f()$ to procedure $g()$ is a tail call if the only thing $f()$ does, after $g()$ returns to it, is itself return. The call is tail-recursive if $f()$ and $g()$ are the same procedure.
- **Partial Redundancy Elimination**
 - Partial-redundancy elimination inserts and deletes computations in the flowgraph in such a way that after the transformation each path contains no more-and, generally, fewer-occurrences of any such computation than before.
- **Register Allocation**
 - Register allocation determines which of the values (variables, temporaries, and large constants) that might profitably be in a machine's registers should be in registers at each point in the execution of a program.

Intermediate Representations

- System of structures that collect essential information about program convenient for analysis, transformation and target code generation
- Two kinds
 - High level IR (platform independent)
 - Low level IR (platform dependent)
- E.g.
 - Abstract Syntax Tree (AST)
 - Static Single Assignment (SSA)
 - gcc: GIMPLE, RTL
 - Clang: LLVM
 - V8 Java Script engine: Hydrogen, Litium

Optimizations for sequential uniprocessor


Reaching Definitions Problem

Control Flow Graph

Data Flow Analysis

Defs & Uses

- **Definition** of the variable is a point in program code where it is assigned to.
- **Use** of variable is a point in code where its value is used.
- Du- and ud-chains are a sparse representation of data-flow information about variables.
- **Du-chain** for a variable connects a definition of that variable to all the uses it may flow to.
- **Ud-chain** connects a use to all the definitions that may flow to it.



Reaching Definitions

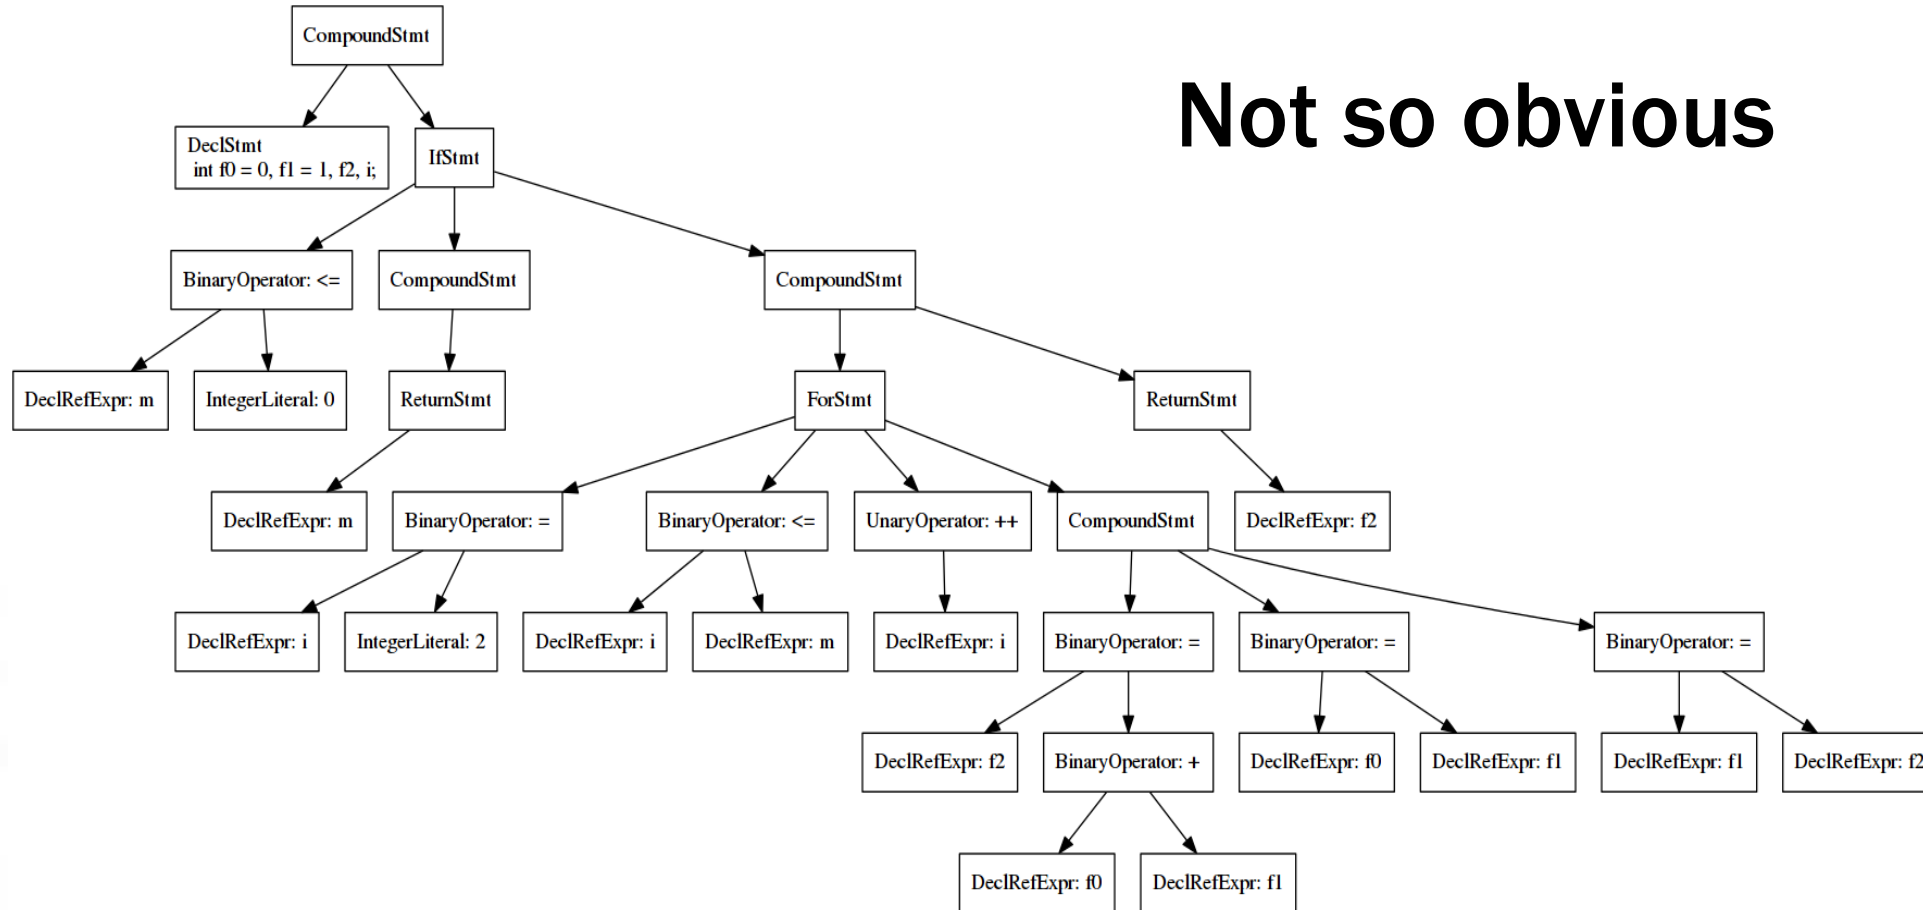
```
int fibonacci(int m)
{
    int f0 = 0, f1 = 1, f2, i;

    if(m <= 1)
    {
        return m;
    }
    else {
        for(i = 2 ; i <= m; ++i)
        {
            f2 = f0 + f1;
            f0 = f1;
            f1 = f2;
        }
        return f2;
    }
}
```

**Whether f2 in
return statement
has definite value?**

Abstract Syntax Tree

Not so obvious



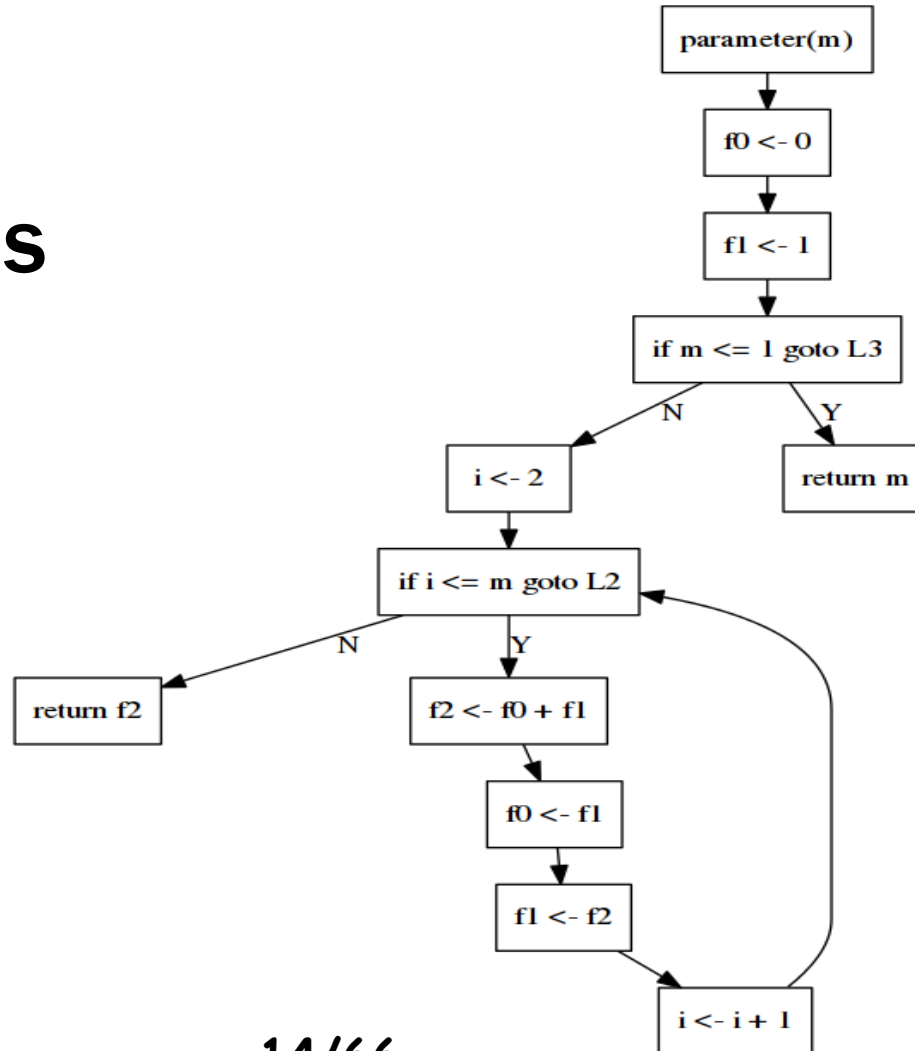
Backend Intermediate Representation

```
1      parameter(m)
2      f0 <- 0
3      f1 <- 1
4      if m <= 1 goto L3
5      i <- 2
6  L1:  if i <= m goto L2
7      return f2
8  L2:  f2 <- f0 + f1
9      f0 <- f1
10     f1 <- f2
11     i <- i + 1
12     goto L1
13  L3: return m
```

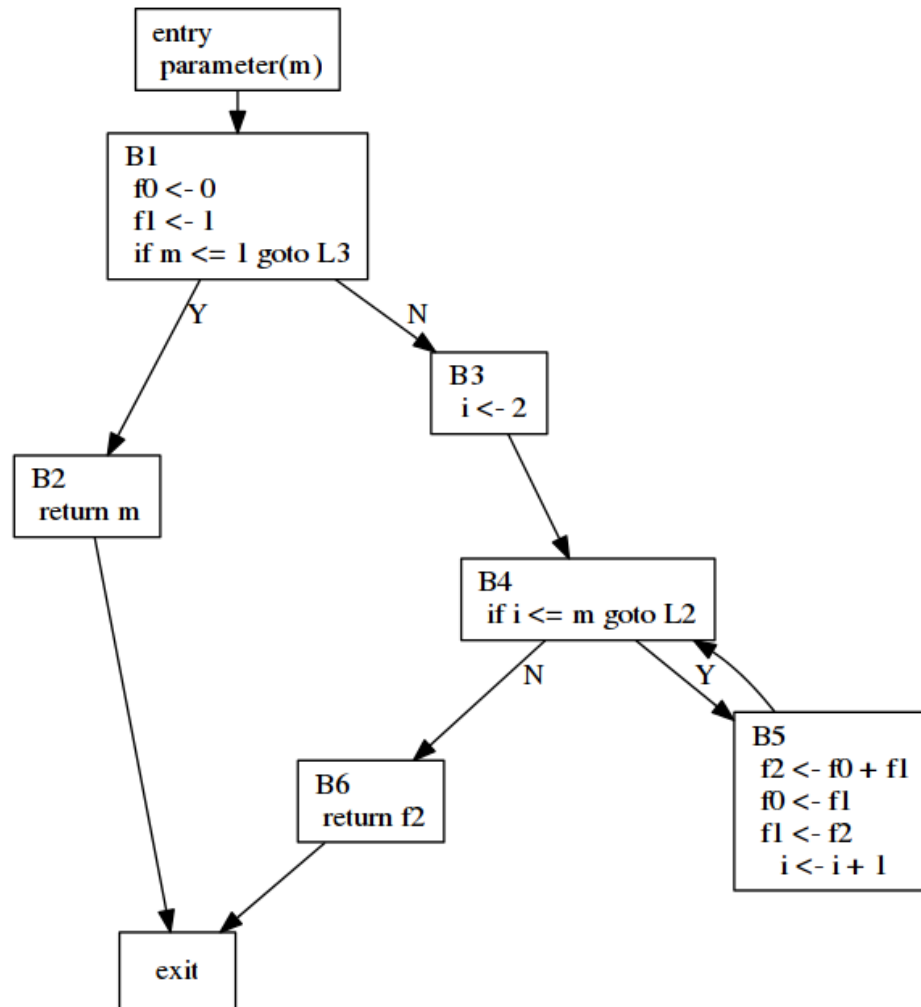
**Can perform
Symbolic
Execution**

Control Flow Chart

Too many nodes
and edges



Control Flow Graph (CFG)



CFG summarizes
control flow with
Basic Blocks and
control edges

Basic Block: a straight-line
code sequence with no branches
in except to the entry and no
branches out except at the exit

Data Flow Analysis: Reaching Definitions

Bit Position	Definition	Basic Block
1	m in line 1	B1
2	f0 in line 2	
3	f1 in line 3	
4	i in line 5	B3
5	f2 in line 8	B5
6	f0 in line 9	
7	f1 in line 10	
8	i in line 11	

```
1      parameter(m)
2      f0 <- 0
3      f1 <- 1
4      if m <= 1 goto L3
5      i <- 2
6  L1:  if i <= m goto L2
7      return f2
8  L2:  f2 <- f0 + f1
9      f0 <- f1
10     f1 <- f2
11     i <- i + 1
12     goto L1
13  L3:  return m
```

Data Flow Analysis: Reaching Definitions

- RCHin – definitions reaching block entry
 - RCHout – definitions reaching block exit
 - PRSV – definitions preserved by block
 - GEN – definitions generated by block
-
- Summarize PRSV and GEN sets for each Basic Block
 - Iteratively distribute information through control edges

Reaching Definitions: Initialization

- $RCHin(entry) = \langle 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0 \rangle$
- $RCHin(i) = \langle 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0 \rangle$ for $\forall i$
- $PRSV(B1) = \langle 0\ 0\ 0\ 1\ 1\ 0\ 0\ 1 \rangle$
- $PRSV(B3) = \langle 1\ 1\ 1\ 0\ 1\ 1\ 1\ 0 \rangle$
- $PRSV(B5) = \langle 1\ 0\ 0\ 0\ 0\ 0\ 0\ 0 \rangle$
- $PRSV(i) = \langle 1\ 1\ 1\ 1\ 1\ 1\ 1\ 1 \rangle$ for other i
- $GEN(B1) = \langle 1\ 1\ 1\ 0\ 0\ 0\ 0\ 0 \rangle$
- $GEN(B3) = \langle 0\ 0\ 0\ 1\ 0\ 0\ 0\ 0 \rangle$
- $GEN(B5) = \langle 0\ 0\ 0\ 0\ 1\ 1\ 1\ 1 \rangle$
- $GEN(i) = \langle 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0 \rangle$ for other i
- $RCHout(i) = \langle 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0 \rangle$ for $\forall i$

Bit Position	Definition	Basic Block
1	m in line 1	B1
2	f0 in line 2	
3	f1 in line 3	
4	i in line 5	B3
5	f2 in line 8	B5
6	f0 in line 9	
7	f1 in line 10	
8	i in line 11	

Reaching Definitions: Equations

$$RCHout(i) = GEN(i) \cup (RCHin(i) \cap PRSV(i)), \forall i$$

$$RCHin(i) = \bigcup_{j \in Pred(i)} RCHout(j), \forall i$$

Reaching Definitions: Result

RCHout(entry) = <0 0 0 0 0 0 0 0>

RCHout(B1) = <1 1 1 0 0 0 0 0>

RCHout(B2) = <1 1 1 0 0 0 0 0>

RCHout(B3) = <1 1 1 1 0 0 0 0>

RCHout(B4) = <1 1 1 1 1 1 1 1>

RCHout(B5) = <1 0 0 0 1 1 1 1>

RCHout(B6) = <1 1 1 1 1 1 1 1>

RCHout(exit) = <1 1 1 1 1 1 1 1>

RCHin(entry) = <0 0 0 0 0 0 0 0>

RCHin(B1) = <0 0 0 0 0 0 0 0>

RCHin(B2) = <1 1 1 0 0 0 0 0>

RCHin(B3) = <1 1 1 0 0 0 0 0>

RCHin(B4) = <1 1 1 1 1 1 1 1>

RCHin(B5) = <1 1 1 1 1 1 1 1>

RCHin(B6) = <1 1 1 1 1 1 1 1>

RCHin(exit) = <1 1 1 1 1 1 1 1>

f2 in block B6 has definite value


Data Flow Analysis: Applications

Problems for DFA:

- Reaching definitions
- Upward exposed uses
- Live variables
- Available expressions
- Copy propagation
- Constant propagation

Variations of DFA

- Forward analysis: from entry to exit
- Backward analysis: from exit to entry
- Confluence operator: union vs intersection



Static Single Assignment Form

Conditional Constant Propagation
Building SSA


Conditional Constant Propagation

```
x = 1;  
y = x + 2;  
if(y > x)  
    y = 5;  
// ...  
... y ...
```



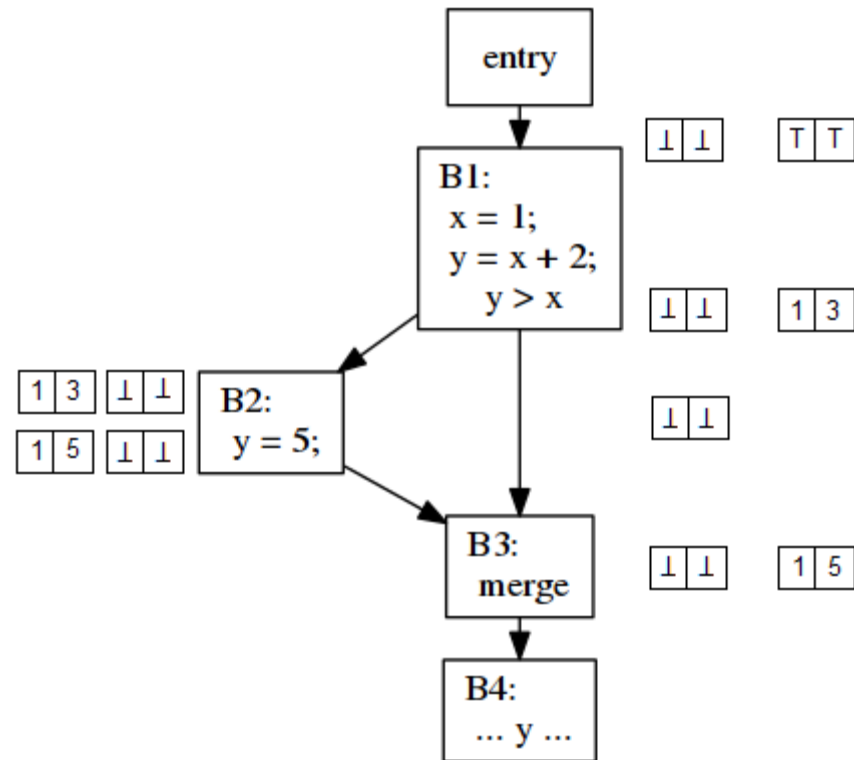
```
x = 1;  
y = 3;           // dead code  
if(true)  
    y = 5;  
// ...  
... 5 ...
```

Conditional Constant Propagation: Lattice



- Lattice for Constant Propagation
 - \perp – yet to be determined
 - T – can't determine / definitely not a constant
- Operators:
 - join(a, b): lowest value above both a and b (join(T, 0) = T, join(0, \perp) = T)
 - meet(a, b): highest value below both a and b (meet(0, \perp) = \perp , meet(T, 1) = 1)
- Evaluation of expressions EVAL(e, Vin):
 - If any argument of e in Vin is T (or \perp) returns T (or \perp), otherwise evaluate as usual


Conditional Constant Propagation: Result



Algorithm Complexity

- Height of lattice == 2 \Rightarrow each vector can change value $2 * V$ times
- Maximal number of iterations: $2 * V * E$ times
- Cost of each iteration: $O(V)$
- Overall algorithm takes $O(EV^2)$ time
- A lot of useless coping unchanged values from one vector to another

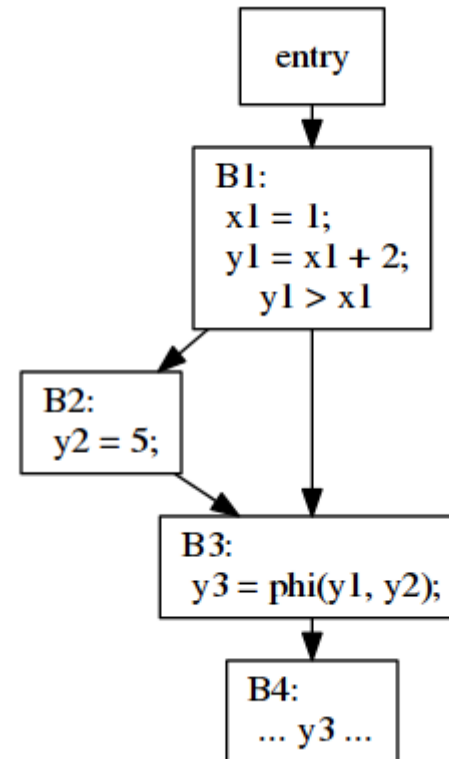
Conditional Constant Propagation: Def-use



- Sparse representation: def-use chains
- Graph for each variable connecting definition to all reachable uses
- Complexity: $O(N^2V)$
- **Loss of accuracy**: does not take into account control flow!

Conditional Constant Propagation: SSA

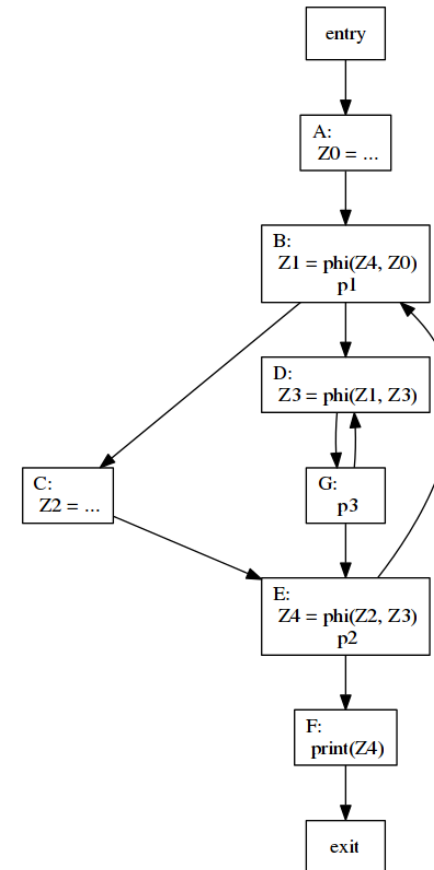

- SSA – Static Single Assignment form
- Each definition creates new version of variable
- ϕ function – pseudo-operation that combines different definitions at merge point to new version
- SSA is a factorized def-use chain which respects control flow
- Constant Propagation in SSA form has complexity $O(EV)$



Building SSA

- Intermediate representation of program in which **every use** of a variable is reached by **exactly one definition**.
- Most program do not satisfy this condition. One requires to insert dummy assignments called ϕ -**functions** to merge multiple definitions.
- Simple algorithm:
 - Insert ϕ -functions for all variables at all CFG merge points
 - Solve reaching definitions
 - Rename all real and dummy assignments of variables uniquely



Building SSA



Building SSA

- In the previous slide, dummy assignment Z3 is not really needed since there is no actual assignment to Z in nodes D and G of the original program
- Minimal SSA Form
 - Place ϕ -functions only where they are really needed
 - Perform direct renaming

Building Minimal SSA



Building Minimal SSA

- Node a **dominates** node b iff every path from entry to b goes through a
- Node u is in **dominance frontier** of node w if w
 - dominates a CFG predecessor v of u , but
 - does not strictly dominate u
- Dominance frontier == control dependence in reverse graph!
- **Iterated dominance frontier**: irreflexive transitive closure of dominance frontier relation
- Where to place ϕ -functions for a variable Z :
 - Let $\text{Assignment} = \{\text{START}\} \cup \{\text{nodes with assignment to } Z\}$
 - Find $I = \text{iterated dominance frontier of Assignment}$
 - Place ϕ -functions in nodes of set I

Optimizations of Parallel Computations

Parallelization overview

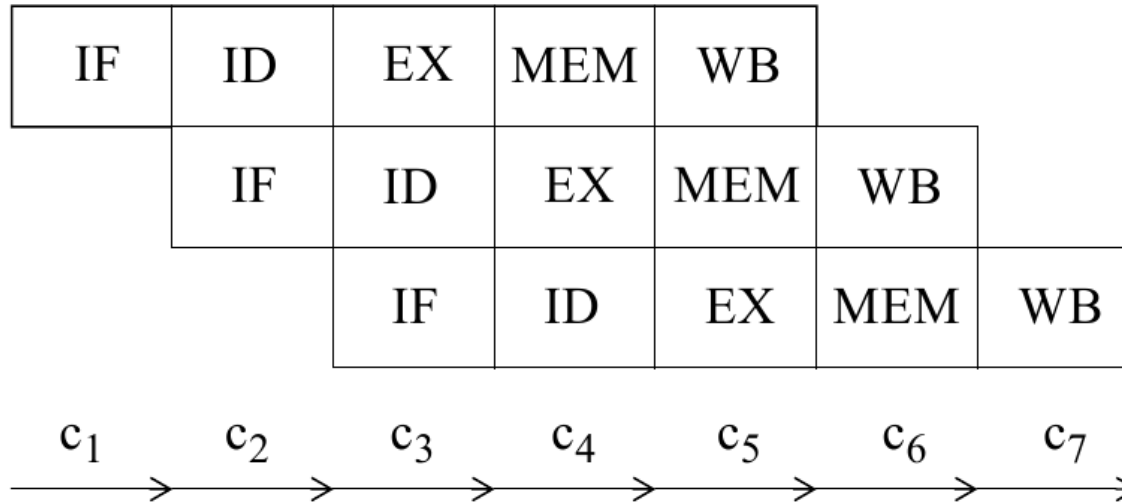
Matrix multiplication case study

Introduction into dependency analysis

Types of Parallelism

- Bit Level Parallelism
 - 8 => 16 => 32 => 64 => ...
- Instruction Level Parallelism
 - Pipelining
 - Superscalar and Very Large Instruction Word (VLIW, EPIC)
- Data Level Parallelism
 - Vector Instructions
 - Graphic Processing Unit (GPU)
- Task Level Parallelism
 - Multicore and Manycore
 - Simultaneous multithreading, Hyperthreading
 - Clusters and Datacenters

Pipelined Instruction Unit




IF – Instruction Fetch
ID – Instruction Decode
EX – Execute
MEM – Memory Access
WB – Write Back

Compiling for Instruction Pipeline

- **Structural hazards** – hardware resources do not support all possible combinations of instruction overlap
- **Data hazards** – result produced by one instruction is needed by subsequent instruction
- **Control hazards** – pipeline stall due to branching instruction

Instruction Pipeline: Structural Hazard



Instruction Pipeline: Data Hazard

LW R1, 0(R2)

ADD R3, R1, R4



IF	ID	EX	MEM	WB		
	IF	ID	stall	EX	MEM	WB

Instruction Pipeline: Control Hazard

Conditional Branch

IF	ID	EX	MEM	WB					
	IF	ID	EX	IF	ID	EX	MEM	WB	
	stall	stall	stall						

Pipelined Execution Unit: Floating Adder




Pipelined Execution Unit: Data Hazard

ADDD R3 , R1 , R2

ADDD R3 , R3 , R4

IF	ID	EX1	EX2	MEM	WB		
	IF	ID	stall	EX1	EX2	MEM	WB

Vector Instruction Unit




- Vector ALU
- Vector registers
- Vector instructions
 - VLOAD VR1, M1
 - VLOAD VR2, M2
 - VADD VR3, VR1, VR2
- Intel: MMX, SSE, AVX
ARM: NEON
PowerPC: AltiVec
Sparc: VIS
MIPS: MSA

Compiling for Vector Processor

- Vector instruction sequence:
VLOAD VR1, A
VLOAD VR2, B
VADD VR3, VR1, VR2
VSTORE VR3, C
- In language supporting vector instructions (Cilk):
 $C[0:63] = A[0:63] + B[0:63]$
- In usual language:
for(i = 0; i < 64; ++i)
 C[i] = A[i] + B[i]
- A slightly modified example:
for(i = 0; i < 64; ++i)
 A[i + 1] = A[i] + B[i]
- In language supporting vector instructions:
 $A[1:64] = A[0:63] + B[0:63]$ **// NOT VECTORIZED**

Multiple-Issue Instruction Unit



- Multiple functional units
 - Integer ALUs
 - Floating ALUs
 - Address ALUs
- Multiple-issue decoder
- **Superscalar**: look ahead in instruction stream for ready-to-execute instructions
- **VLIW**: execute single wide instruction in each cycle

More than one instruction per cycle

Scheduling for Pipelined Processor

```
LD    R1, A
LD    R2, B
FADD  R3, R1, R2
ST    R3, X
LD    R4, C
FADD  R5, R3, R4
ST    R5, Y
```

7 instruction + 4 delays = 11 cycles

```
LD    R1, A
LD    R2, B
LD    R4, C
FADD  R3, R1, R2
FADD  R5, R3, R4
ST    R3, X
ST    R5, Y
```

7 instruction + 1 delay = 8 cycles

Scheduling for Multiple-issue Processor

Suppose we have processor with 2 loads but 1 addition per cycle

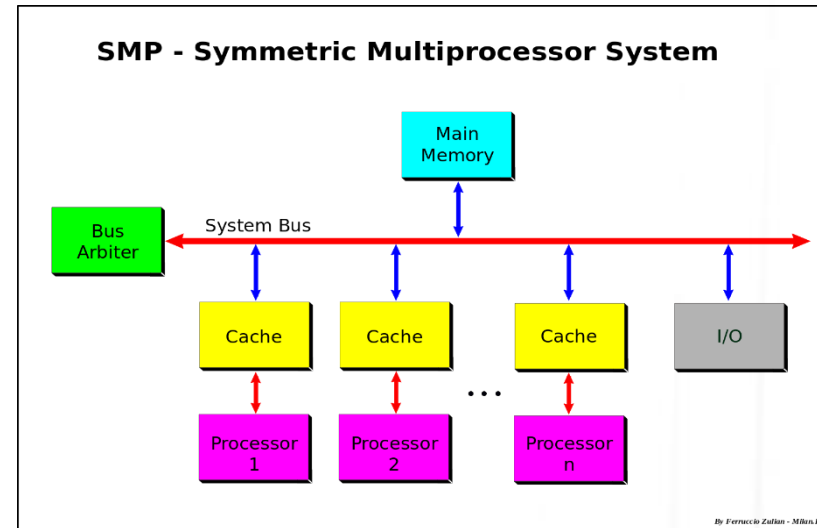
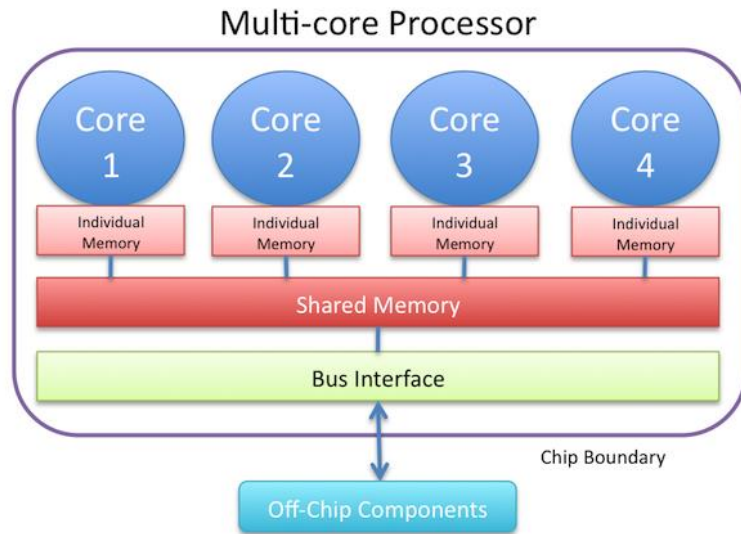
```
LD    R1, A
LD    R2, B
FADD  R3, R1, R2
ST    R3, X
LD    R4, C
LD    R5, D
FADD  R6, R4, R5
ST    R6, Y
```

LD R1, A	LD R4, C
LD R2, B	LD R5, D
delay	delay
FADD R3, R1, R2	delay
STD R3, X	FADD R6, R4, R5
empty	ST R6, Y

VLIW – compiler generates wide instructions with two slots

Superscalar – compiler schedules instructions to fit in look-ahead window

Processor Parallelism



- **Multicore processor:** single component with two or more independent actual processing unit
- **Multiprocessor system:** two or more CPUs within a single computer system

Compiling for Processor Parallelism

```
#pragma omp parallel for  
for(i = 0; i < N; ++i)  
    C[i] = A[i] + B[i];
```

Parallelized / Vectorized

```
#pragma omp parallel for  
for(i = 0; i < N; ++i)  
    A[i - 1] = A[i] + B[i];
```

Non-Parallelized / Vectorized

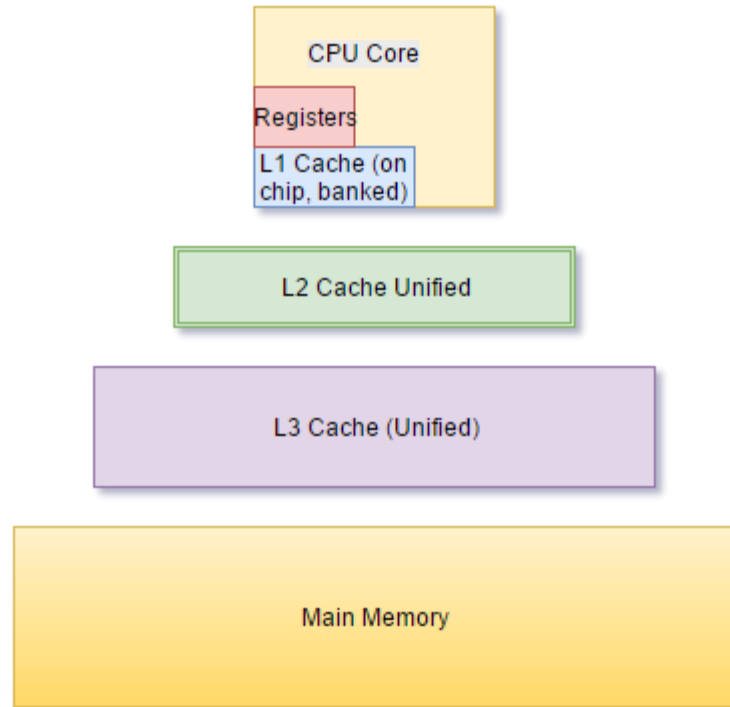
```
#pragma omp parallel for  
for(i = 0; i < N; ++i)  
    A[i + 1] = A[i] + B[i];
```

Non-Parallelized / Non-Vectorized

```
#pragma omp parallel for  
for(i = 0; i < N; ++i)  
    S = A[i] + B[i];
```

Non-Parallelized / Non-Vectorized

Memory Hierarchy



- **Latency** – number of cycles required to deliver single data element from memory
- **Bandwidth** – number of elements delivered on each cycle
- **Latency avoidance** – reduce latency in computations (caches)
- **Latency tolerance** – doing something else while data is being fetched

Compiling for Memory Hierarchy

```
for(i = 0; i < N; ++i)
  for(j = 0; j < M; ++j)
    A[i] = A[i] + B[j];
```

```
for(jj = 0; jj < M; jj += L)
  for(i = 0; i < N; ++i)
    for(j = jj; j < jj + L; ++j)
      A[i] = A[i] + B[j];
```

- Access for array A fits into cache
- Access for array B misses for large M
- Suppose cache size larger than L array elements
- **Strip-mining** loop

Matrix Multiplication: Scalar Uniprocessor

```
for(i = 0; i < N; ++i)
  for(j = 0; j < N; ++j)
  {
    C[i][j] = 0.0;
    for(k = 0; k < N; ++k)
      C[i][j] += A[i][k] * B[k][j];
  }
```

Straightforward algorithm – good performance on scalar uniprocessor

Matrix Multiplication: Pipelined Processor

```
for(i = 0; i < N; i += 4)
  for(j = 0; j < N; ++j)
  {
    C[i][j]    = 0.0;
    C[i+1][j]  = 0.0;
    C[i+2][j]  = 0.0;
    C[i+3][j]  = 0.0;
    for(k = 0; k < N; ++k)
    {
      C[i][j]    += A[i][k]    * B[k][j];
      C[i+1][j] += A[i+1][k]  * B[k][j];
      C[i+2][j] += A[i+2][k]  * B[k][j];
      C[i+3][j] += A[i+3][k]  * B[k][j];
    }
  }
```

Loop unrolling

Matrix Multiplication: Vector Processor

```
for(i = 0; i < N; i += 64)
  for(j = 0; j < N; ++j)
  {
    C[i:i+63][j] = 0.0;
    for(k = 0; k < N; ++k)
      C[i:i+63][j] += A[i:i+63][k] * B[k][j];
  }
```

Innermost loop is not vectorized,
So vectorize outermost loop

Matrix Multiplication: Superscalar Processor

```
for(i = 0; i < N; i += 4)
  for(j = 0; j < N; j += 4)
  {
    C[i:i+3][j]    = 0.0;
    C[i:i+3][j+1] = 0.0;
    C[i:i+3][j+2] = 0.0;
    C[i:i+3][j+3] = 0.0;
    for(k = 0; k < N; ++k)
    {
      C[i:i+3][j]    += A[i:i+3][k] * B[k][j];
      C[i:i+3][j+1] += A[i:i+3][k] * B[k][j+1];
      C[i:i+3][j+2] += A[i:i+3][k] * B[k][j+2];
      C[i:i+3][j+3] += A[i:i+3][k] * B[k][j+3];
    }
  }
}
```

4-way simultaneous issue, 4 floating-point multiply-adders, 4 pipeline stages

Matrix Multiplication: Parallel Processor

```
#pragma omp parallel for
for(i = 0; i < N; ++i)
  for(j = 0; j < N; ++j)
  {
    C[i][j] = 0.0;
    for(k = 0; k < N; ++k)
      C[i][j] += A[i][k] * B[k][j];
  }
```

Parallelize outermost loop

Matrix Multiplication: Memory Hierarchy

```
for(ii = 0; ii < N; ii += L)
  for(jj = 0; jj < N; jj += L)
  {
    for(i = ii; i < ii + L; ++i)
      for(j = jj; j < jj + L; ++j)
        C[i][j] = 0.0;
    for(kk = 0; kk < N; kk += L)
      for(i = ii; i < ii + L; ++i)
        for(j = jj; j < jj + L; ++j)
          for(k = kk; k < kk + L; ++k)
            C[i][j] += A[i][k] * B[k][j];
  }
```

Cache size greater than $3L^2$ data item size

Introduction into Dependency Analysis

- Manual optimization for particular architecture leads to **non-portable** programs
- All these optimizations should be performed by **compiler**
- But optimizations should keep program's **correctness**
- Formalization of correctness: **dependency** between program statements
- Statement S2 **depends** on statement S1:
 - Both S1 and S2 access **the same memory**
 - There is **control path** from S1 to S2
 - At least one of these accesses is **write**

Loop Nest


```
    for(i1 = L1; i1 < U1; i1 += ST1)
      for(i2 = L2; i2 < U2; i2 += ST2)
        ....
        for(in = Ln; in < Un; in += STn)
        {
S1:      A[f1(i1, i2, ..., in), ..., fm(i1, i2, ..., in)] = ....
S2:      ... = A[g1(i1, i2, ..., in), ..., gm(i1, i2, ..., in)];
        }
```

Normalized loop:

- $L_i == 0$
- $ST_i == 1$


Types of Dependency

```
S1:  pi = 3.141592;  
S2:  r = 5.0;  
S3:  area = pi * r * r;
```



Data dependency


```
S1:  if(x != 0.0)  
S2:    a = a / x;
```



Control dependency


Data Dependency Classification

```
S1:  X = ...  
S2:  ... = X
```




True dependence

```
S1:  ... = X  
S2:  X = ...
```



Anti-dependence

```
S1:  X = ...  
S2:  X = ...
```



Output dependence

Loop-carried and Loop-independent

```
for(i = 0; i < N; ++i)
{
    A[i + 1] = F[i];
    F[i + 1] = A[i];
}
```

Loop-carried dependence

```
for(i = 0; i < N; ++i)
{
    A[i] = ... ;
    ... = A[i];
}
```

Loop-independent dependence

Parallelization and Vectorization

- It is valid to convert a sequential loop to a **parallel loop** if the loop **carries no dependence**.
- A statement contained in at least one loop can be **vectorized** if the statement is **not** included in any **cycle of dependences**.

Polyhedral Optimization Frameworks

- Graphite: Gimple Represented as Polyhedra
- <http://gcc.gnu.org/wiki/Graphite>
- Polly: LLVM Framework for High-Level Loop and Data-
Locality Optimizations
- <http://polly.llvm.org/>

Bibliography

1. Steven Muchnick. Advanced Compiler Design and Implementation, 1997
2. Randy Allen, Ken Kennedy. Optimizing Compilers for Modern Architectures: A Dependence-based Approach, 2001.
3. D. Loveman, R. Faneuf. Program Optimization – Theory and Practice. Massachusetts Computer Associates Inc.
4. A. Appel. SSA is Functional Programming.
5. R. Cytron & al. Efficiently Computing Static Single Assignment Form and the Control Dependence Graph.
6. C. McConnell, R. Johnson. Using SSA Form in a Code Optimizer.
7. R. Kennedy & al. Partial Redundancy Elimination in SSA Form
8. K. Cooper & al. Operator Strength Reduction



Thank you!