

Compiler Construction: Practical Introduction

System Course for SRR Engineers

Samsung Research Russia
Moscow 2019

Lecture 9

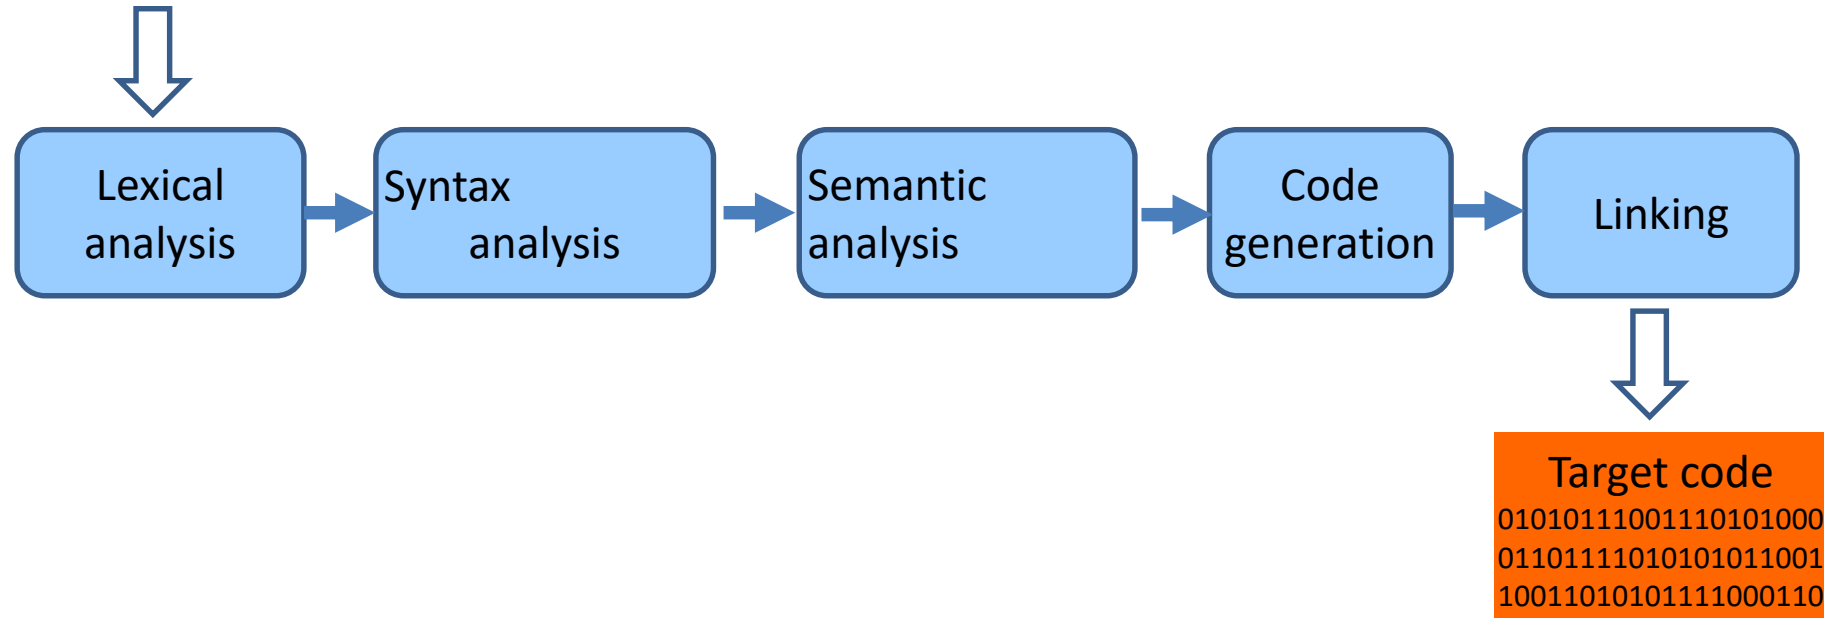
Modern Compiler Architecture

- Conventional monolithic compiler
- LLVM compiler infrastructure
- V8 JavaScript Engine
- OpenCL framework

Compilation: An Ideal Picture

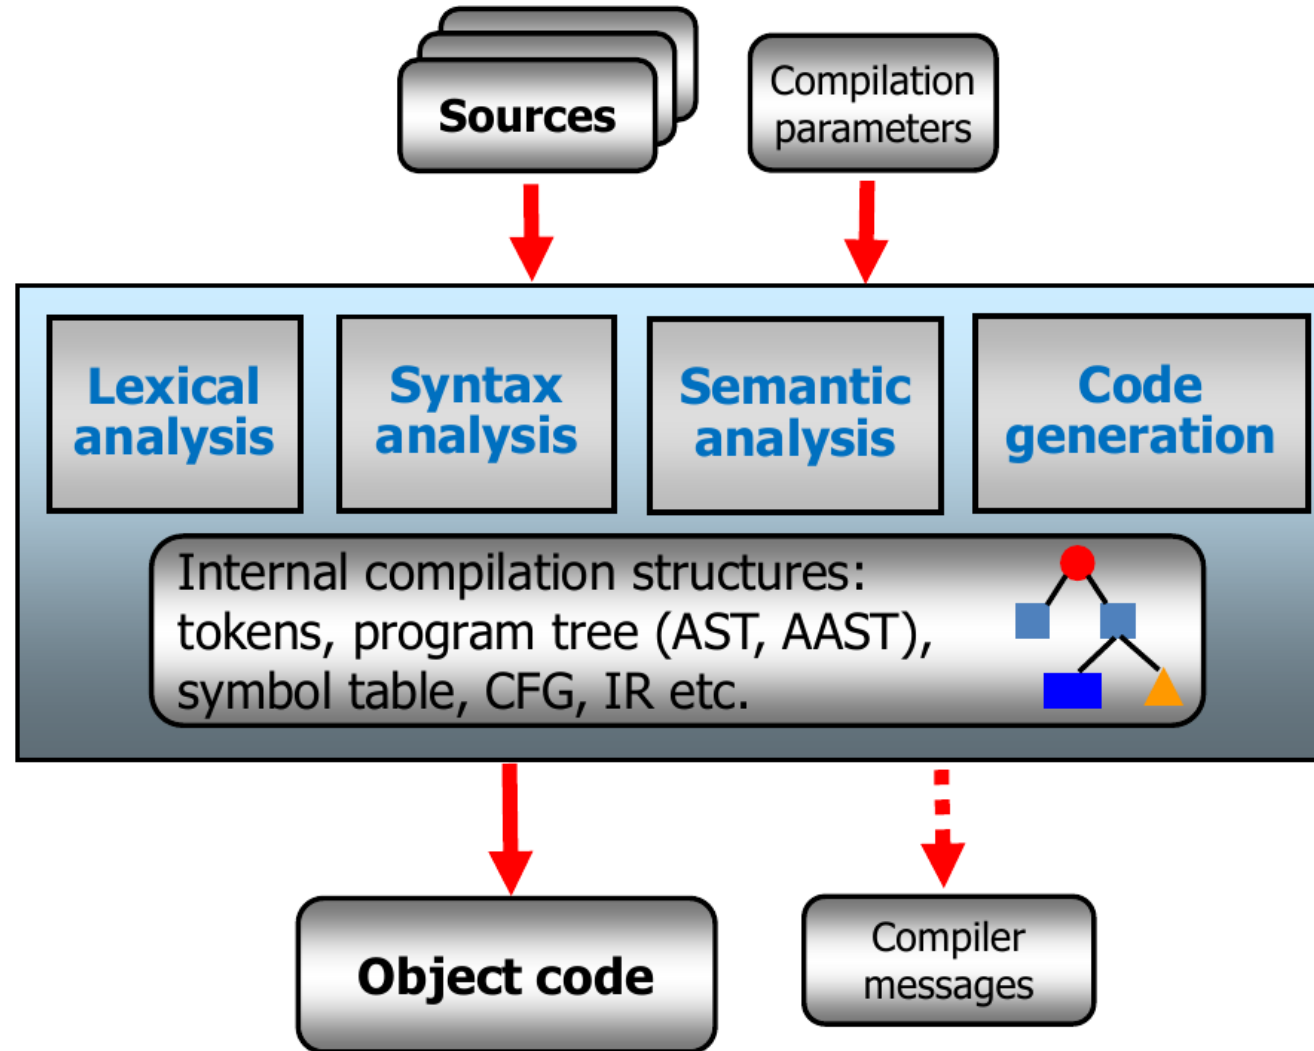
*A program written by a human
(or by another program)*

Source program
text



*A program binary image
suitable for immediate
execution by a machine*

Conventional Compiler



Toward Multi-Platform Compiler

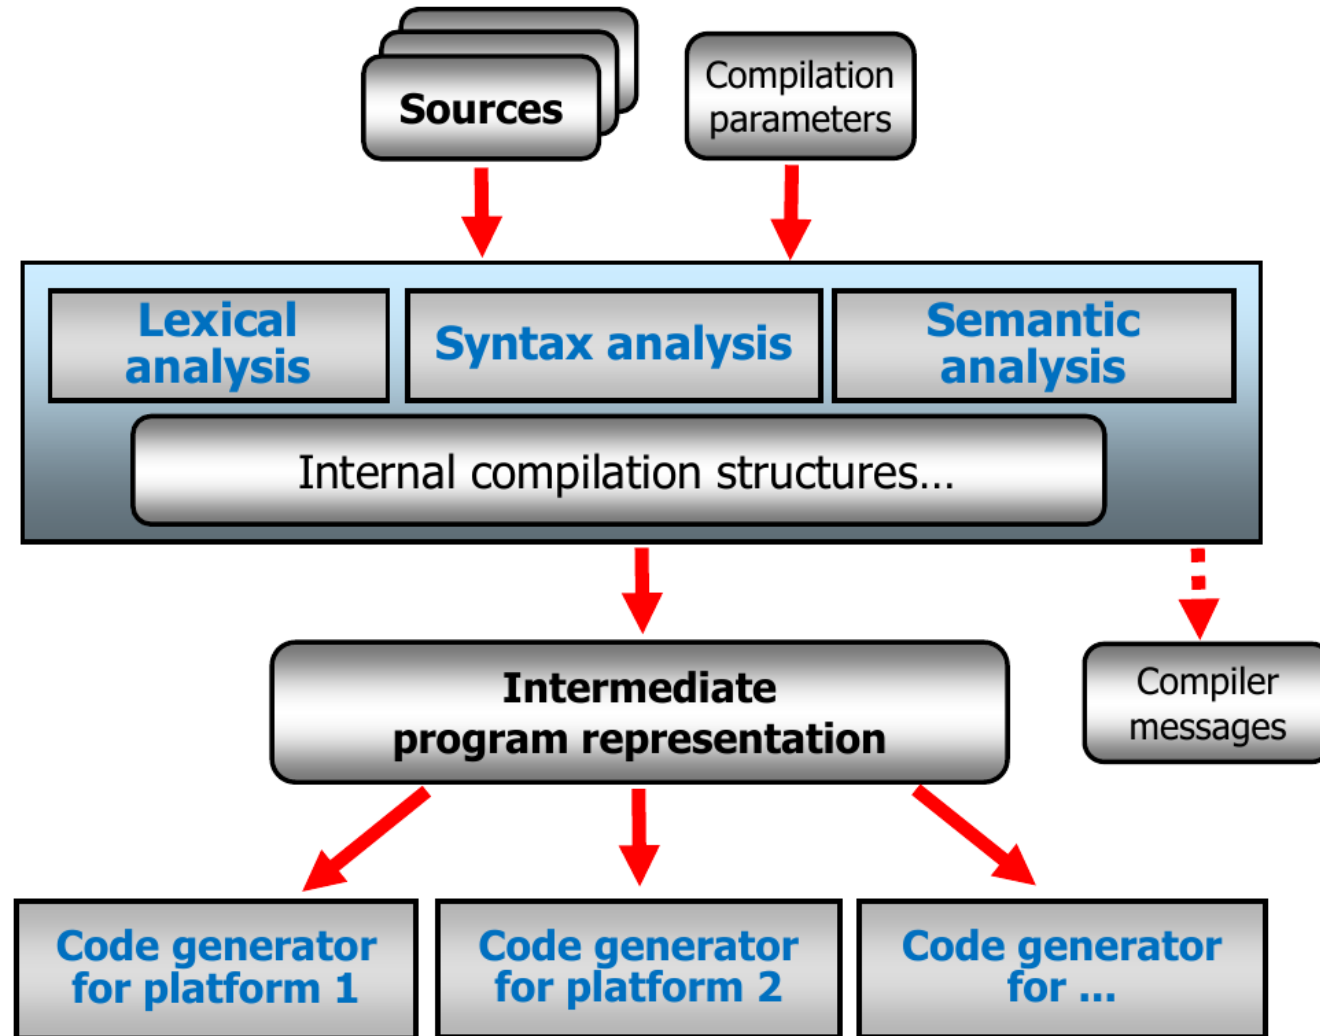
Issues with monolithic compiler

- One needs to rewrite compiler from scratch for every new programming language
- One needs to rewrite compiler from scratch for every new processor architecture
- Compiler optimizations are not re-usable.

Multi-language multi-platform compiler

- Design common Intermediate Representation
- All source programming languages compile into the same IR
- Generate code to all processors from the same IR

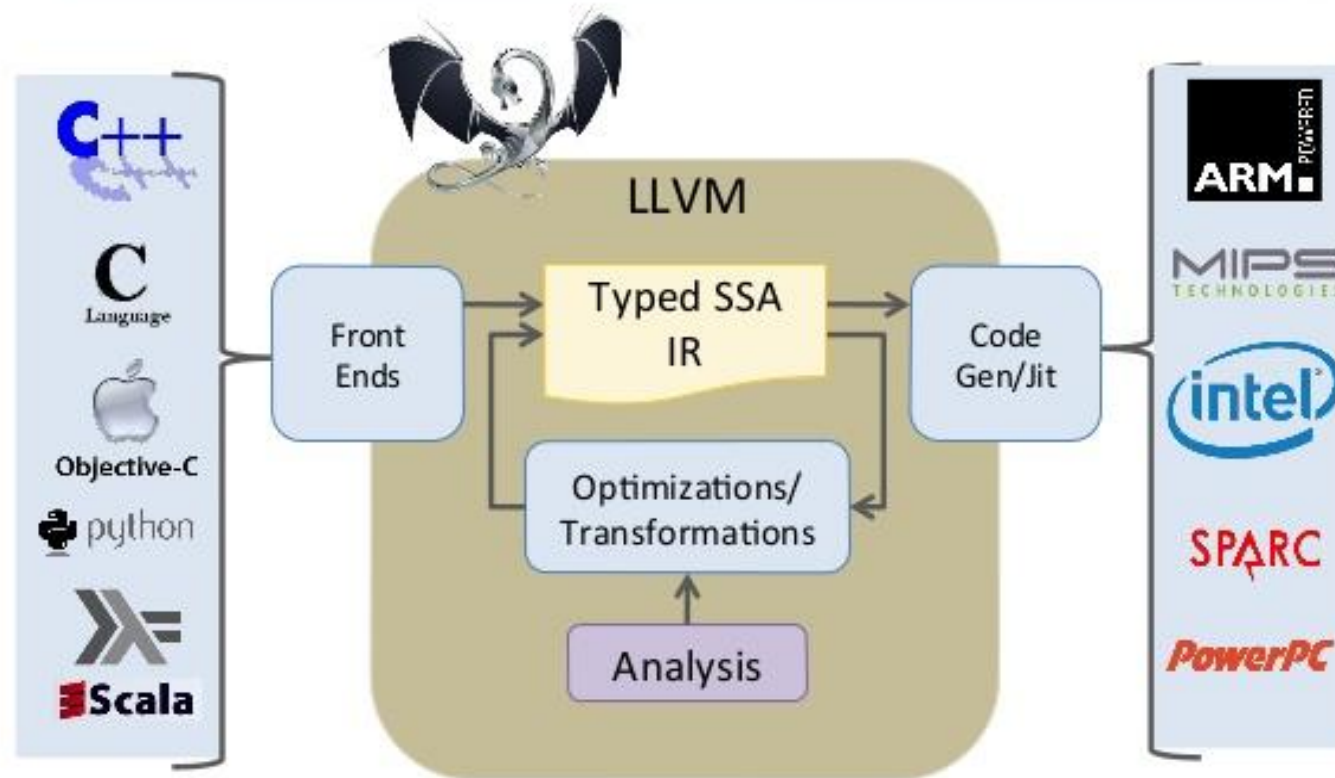
Multi-platform Architecture



LLVM Infrastructure

LLVM Compiler Infrastructure

[Lattner et al.]



LLVM Primary Components

- **The LLVM Virtual Instructions Set:**
 - The common language- and target-independent IR.
 - Internal (IR) and external (persistence) representations.
- **The collection of well-integrated libraries:**
 - Analysis, optimizations, code generators, JIT compiler, garbage collection support, profiling, ...
- **The collection of tools built from the libraries:**
 - Assemblers, automatic debugger, linker, code generators, compiler driver, modular optimizer, ...

LLVM IR: Goals

- Easy to produce, understand, and define.
- Language- and Target-Independent:
 - AST-level IR is not very feasible
 - Every analysis/xform must know about 'all' languages.
- One IR for analysis and optimizations:
 - IR must be able to support aggressive IPO, loop opts, scalar opts, ... high- and low-level optimizations.
- Optimize as much as early as possible:
 - Can't postpone everything until link- or runtime.
 - No lowering in the IR.

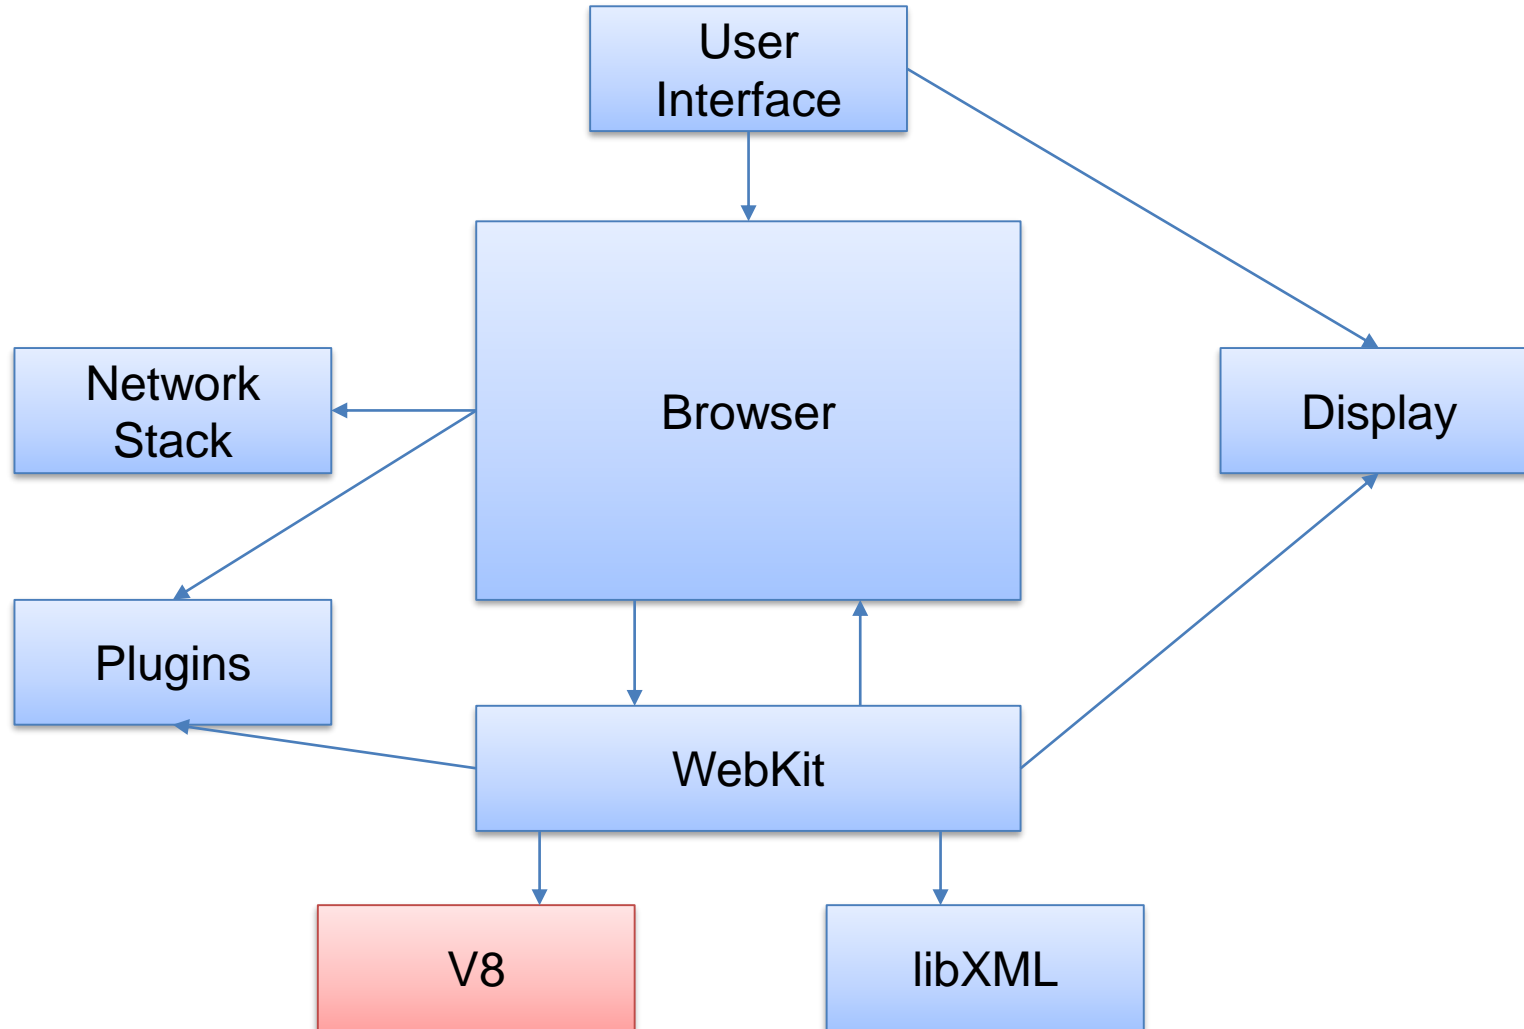
Five Points of LLVM IR

- **Extremely simple IR to learn and use:**
 - 1-to-1 correspondence between .ll, .bc, and C++ IR.
- **Powerful and modular optimizer:**
 - Easy to extend, or just use what is already there.
- **Clean and modular code generator:**
 - Easy to retarget, easy to replace/tweak components.
- **Many “productivity tools” (bugpoint, verifier).**
- **Active development community, good documentation.**

V8 JavaScript Engine

- JavaScript language is the only programming language for the Web (Webassembly is still not in production)
 - Invented by Brendan Eich in 1995, Netscape Navigator
 - Prototype based objects
 - Dynamic typing
 - Functions as a first class citizen
- Google V8 is the one of the fastest JavaScript Engines
 - Part of the Google Chrome web browser
 - Part of the Node.js server
 - Just-in-Time compilation

V8 in Chrome



V8 in Node.js

Node standard library

Node bindings
(socket, http, &c)

V8

Thread
pool

Event
loop

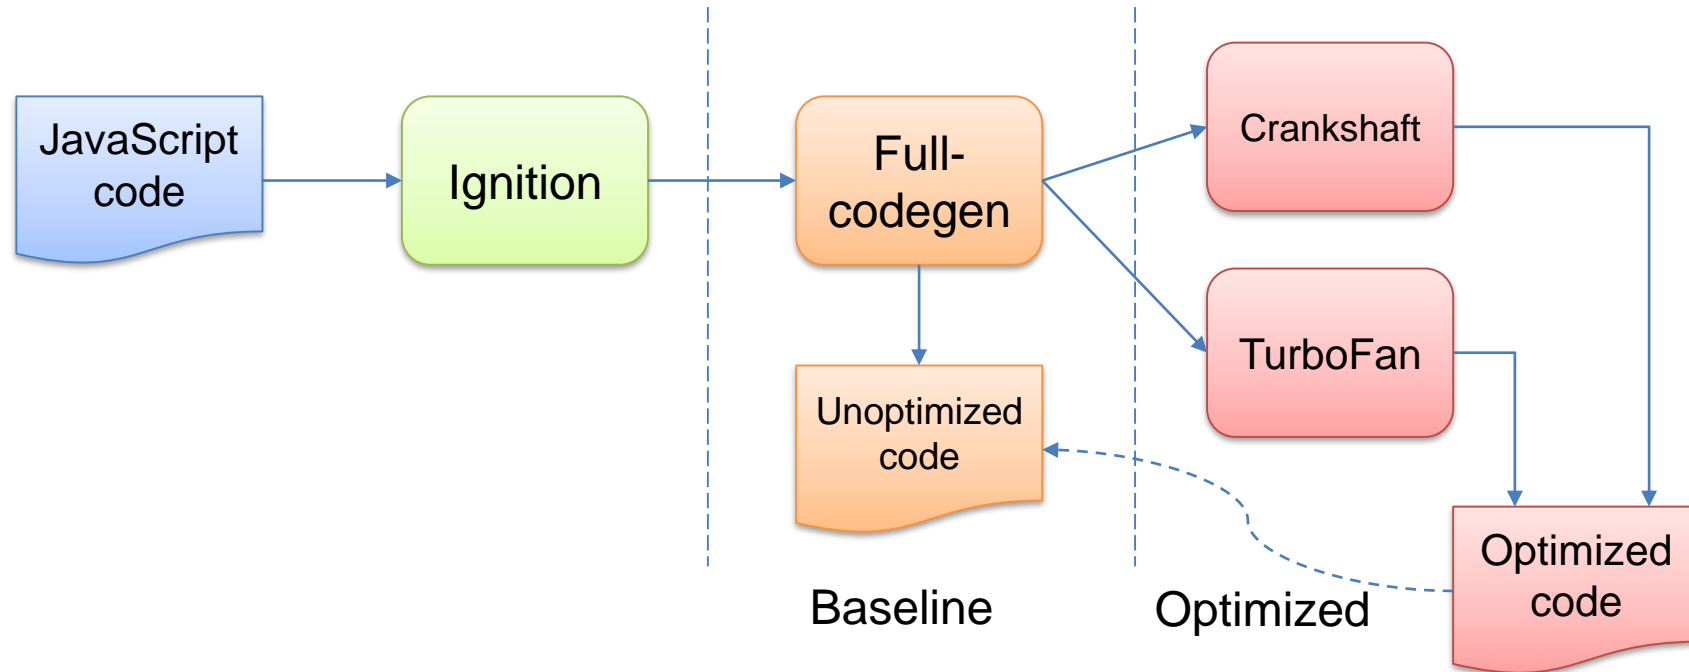
DNS

crypto

JavaScript Compilation Issues

- Dynamic typing
 - Any variable can be assigned object of any type
 - The type of variable can be changed during program execution
- Prototype based objects
 - Object are created by cloning other objects
 - Every object has link to its prototype object
 - Prototype object can be changed during program execution
- `eval()` function: new code can be created at run time

V8 Architecture



- JavaScript source code compiled to internal representation
- Full-codegen generates generic unoptimized machine code
- Hot regions are compiled to optimized code

V8 Compilation Strategies




By @addyosmani

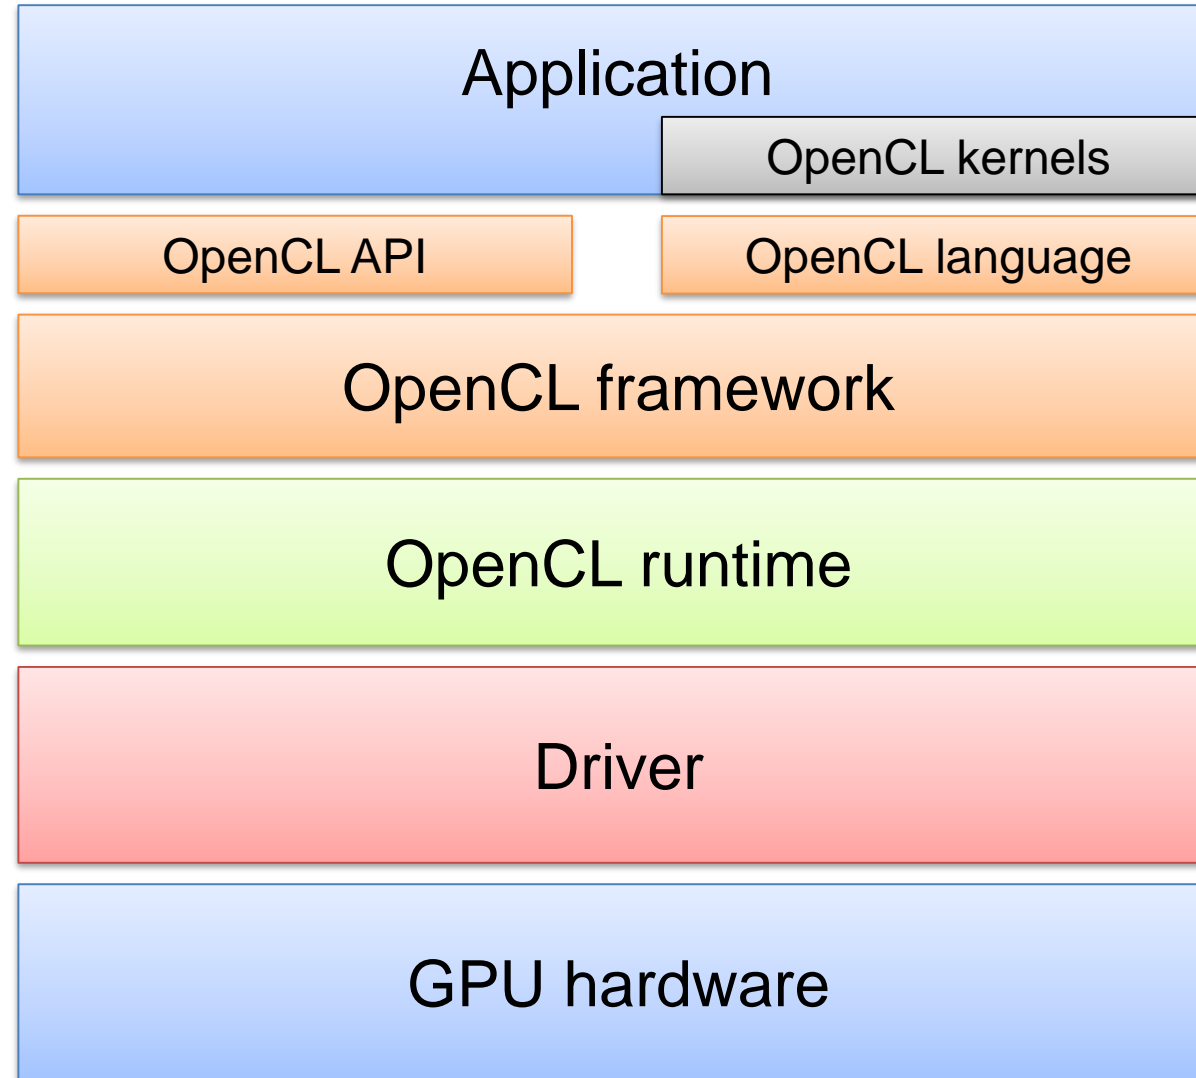
- Fast unoptimized code generation
- Inline Caching: virtual table at invocation site
- Hot region profiling: counters insertion
- Type-dependent optimization of hot-regions
- On-stack Replacement of recompiled code

OpenCL

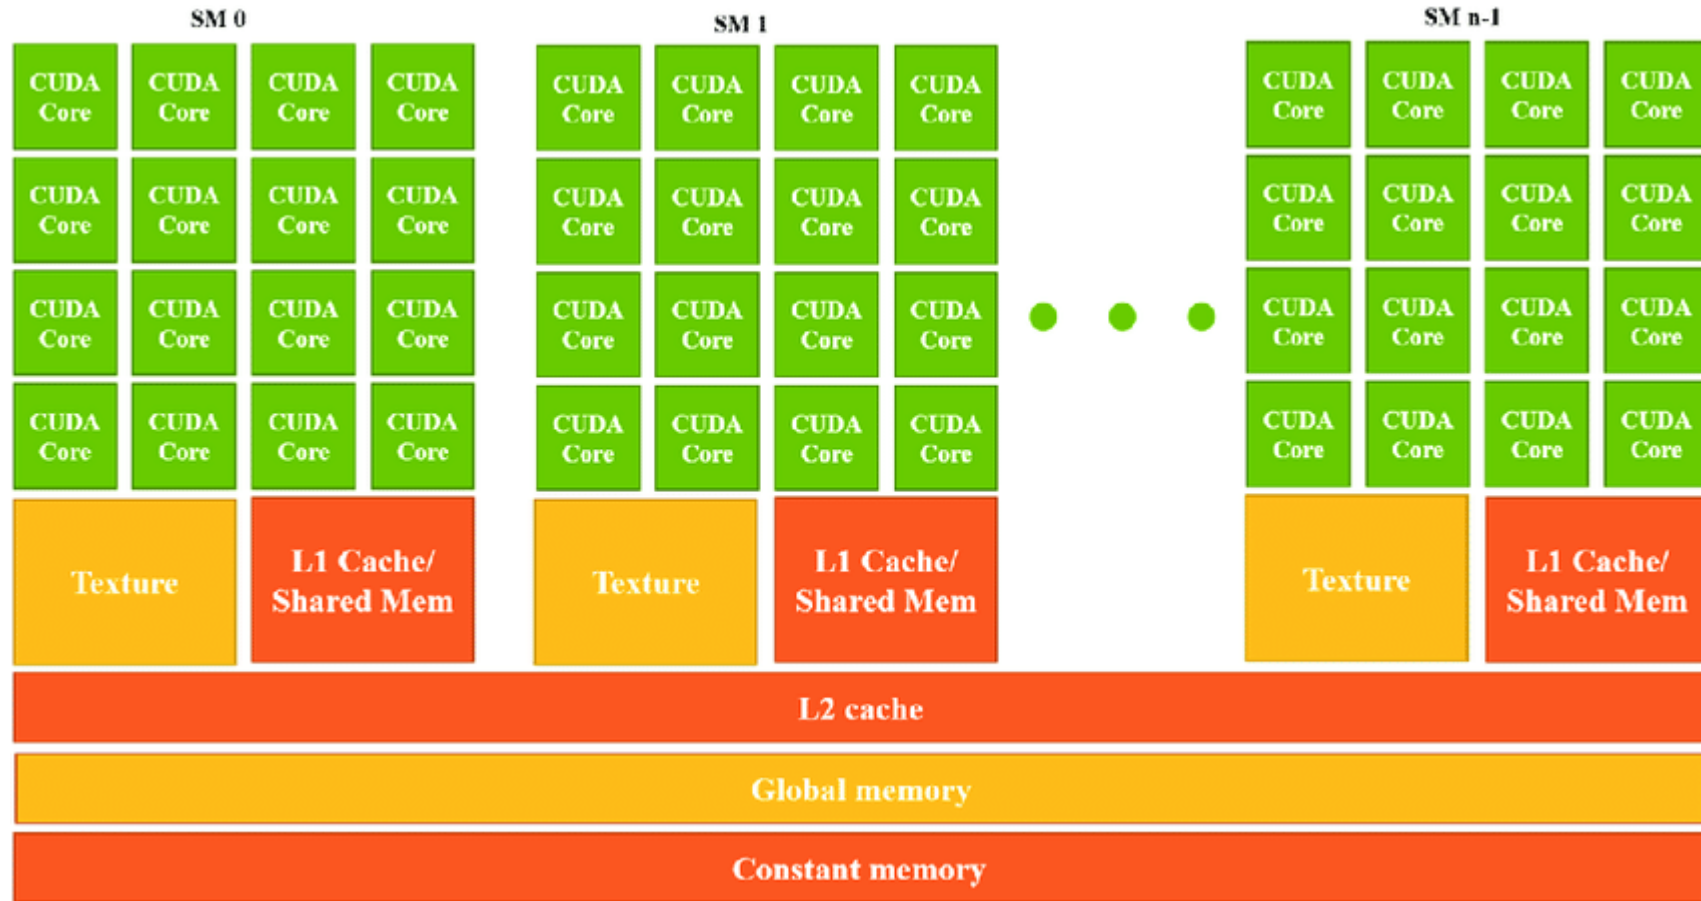
- OpenCL™ (Open Computing Language) is the open standard for cross-platform, parallel programming of diverse processors.
- OpenCL specifies:
 - A programming language based on C99 and C++11
 - Application programming interfaces (APIs) to control the platform and execute programs on the compute devices.
- OpenCL is constituted of:
 - OpenCL programming language compiler
 - OpenCL APIs library
 - Driver to control execution



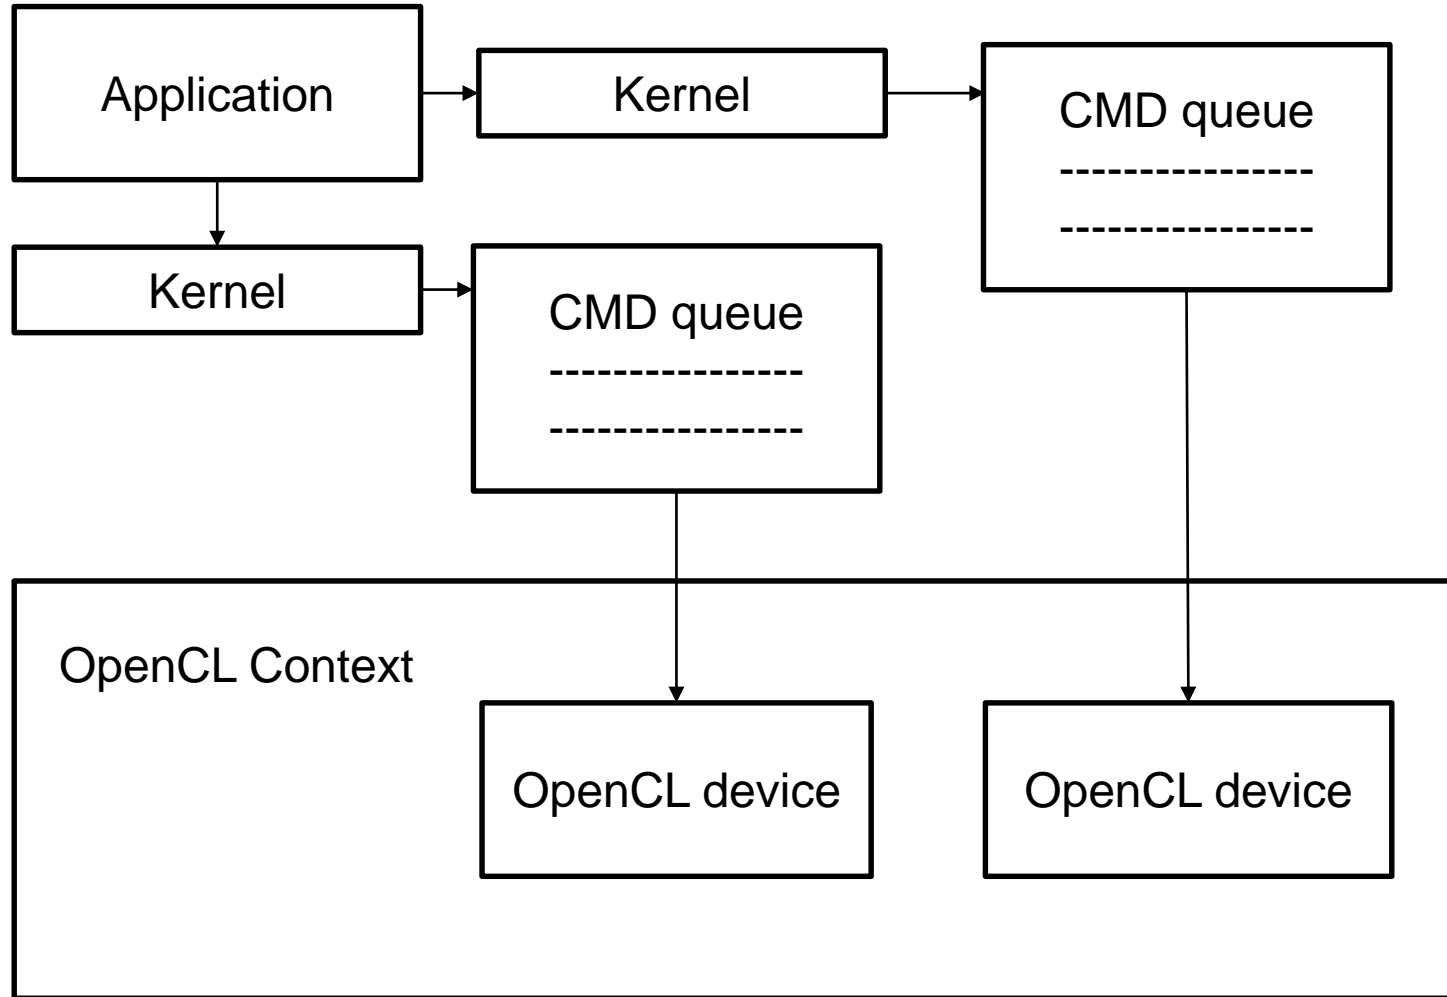
OpenCL Architecture



GPU Architecture



OpenCL Application



OpenCL Program

- An OpenCL «program» contains one or more «kernels» and any supporting routines that run on target device
- An OpenCL kernel is the basic unit of code that can be executed on a target device.

```
// Multiplies A*x, leaving the result in y.  
// A is a row-major matrix, meaning the (i,j) element is at A[i*ncols+j].  
__kernel void matvec(__global const float *A, __global const float *x,  
                    uint ncols, __global float *y)  
{  
    size_t i = get_global_id(0);          // Global id, used as the row index  
    __global float const *a = &A[i*ncols]; // Pointer to the i'th row  
    float sum = 0.f;                       // Accumulator for dot product  
    for (size_t j = 0; j < ncols; j++) {  
        sum += a[j] * x[j];  
    }  
    y[i] = sum;  
}
```

OpenCL Host Program

```
    clGetDeviceIDs(NULL, CL_DEVICE_TYPE_GPU, num, devices, NULL);
    cl_context context = clCreateContextFromType(NULL,
CL_DEVICE_TYPE_GPU, NULL, NULL, NULL);
    clGetDeviceIDs(NULL, CL_DEVICE_TYPE_DEFAULT, 1, devices,
NULL);
    cl_command_queue queue = clCreateCommandQueue(context,
devices[0], 0, NULL);
    cl_mem memobjs[] = { clCreateBuffer(context, CL_MEM_READ_ONLY
| CL_MEM_COPY_HOST_PTR, sizeof(float) * 2 * NUM_ENTRIES, NULL,
NULL),
                                clCreateBuffer(context,
CL_MEM_READ_WRITE, sizeof(float) * 2 * NUM_ENTRIES, NULL, NULL) };
    cl_program program = clCreateProgramWithSource(context, 1, (const
char **)& KernelSource, NULL, NULL);
    clBuildProgram(program, 0, NULL, NULL, NULL, NULL);
    cl_kernel kernel = clCreateKernel(program, "fft1D_1024", NULL);
    clSetKernelArg(kernel, 0, sizeof(cl_mem), (void *)&memobjs[0]);

    clEnqueueNDRangeKernel(queue, kernel, 1, NULL, global_work_size,
local_work_size, 0, NULL, NULL);
```

OpenCL Compiler

- OpenCL compiler works a part of OpenCL framework
- OpenCL compiler is a library called by driver
- It compiles kernels and supporting routines to target specific code
- A driver either schedule execution of compiled kernels on device or save it as a binary
- OpenCL compiler compiles the body of a loop nest
- Loop boudaries are managed by the driver